

# LiViFuser-Nav

Geometry-Aware Dense Vision + Sparse 2D LiDAR Fusion  
for Uncertainty-Aware Local Navigation on Low-Cost Robots

Architecture Specification v1.1  
Final Locked for Pilot Implementation

STATUS: FINAL LOCKED — READY FOR PILOT IMPLEMENTATION

|                 |                                                                          |
|-----------------|--------------------------------------------------------------------------|
| Platform        | TurtleBot3 Burger (differential drive)                                   |
| Sensors         | LDS-03 (2D LiDAR 360°) + Raspberry Pi Camera Module v3                   |
| Vision Backbone | Frozen DINOv3-S+/16                                                      |
| Policy          | Heteroscedastic ACT (H=8) + spatial cross-attention decoder              |
| Uncertainty     | Aleatoric ( $\sigma$ ) + Epistemic/OOD (Mahalanobis on DINO features)    |
| Control         | 10 Hz target   Receding-horizon (1-step)   Independent safety supervisor |
| Date Locked     | 2026-07-27                                                               |

This is the final locked specification. It incorporates geometry-aware fusion with spatial structure preserved to the action head, heteroscedastic uncertainty with variance-collapse protection, free Mahalanobis OOD gating on frozen DINOv3 features, mandatory goal conditioning, receding-horizon execution, and an independent safety supervisor.

# 1. Research Positioning & Contribution

## 1.1 Contribution Statement

A geometry-aware gated fusion policy that combines dense self-supervised visual tokens with full-surround sparse 2D LiDAR, dual uncertainty signals (aleatoric + free Mahalanobis OOD), goal-conditioned temporal memory, and safety-constrained receding-horizon execution for real-time navigation on low-cost differential-drive robots.

## 1.2 Differentiation

| Work           | Focus                                    | Our Differentiation                                                                                                     |
|----------------|------------------------------------------|-------------------------------------------------------------------------------------------------------------------------|
| ViLINT (2026)  | RGB + 3D LiDAR, diffusion, larger robots | Sparse 2D LiDAR + dense SSL vision on small robots; FOV-masked geometry-aware fusion; dual uncertainty; explicit safety |
| NoMaD / ViNT   | Vision-only foundation models            | Structured sparse LiDAR fusion + uncertainty + goal-conditioned temporal memory                                         |
| Classical Nav2 | Hand-engineered local planners           | Fully learned local policy with modern dense features and calibrated uncertainty                                        |

## 1.3 Explicitly Not Claimed

- Gated cross-attention itself is not claimed as novel.
- Aleatoric  $\sigma$  alone is not claimed to detect OOD; Mahalanobis supplies the epistemic signal.
- Contribution is the complete system formulation + empirical validation on low-cost hardware, including the finding that frozen SSL features can serve as a free real-time OOD gate.

# 2. Final Architecture (v1.1)

## 2.1 End-to-End Data Flow



## 2.2 Component Specifications

| Component      | Choice                           | Key Details                                                      |
|----------------|----------------------------------|------------------------------------------------------------------|
| Vision Encoder | Frozen DINOv3-S+/16              | No fine-tuning in v1. Spatial structure preserved.               |
| Visual Tokens  | 32–64 after fixed pooling        | Fixed spatial pooling first (easier to debug).                   |
| LiDAR Encoder  | 1D CNN + angular PE              | Input [r, sin $\theta$ , cos $\theta$ , validity]. 72–90 tokens. |
| Fusion         | Geometry-aware Cross-Attn + Gate | FOV mask from extrinsics. Rear = LiDAR-only path.                |

| Component      | Choice                          | Key Details                                         |
|----------------|---------------------------------|-----------------------------------------------------|
| Temporal path  | Pooled features → GRU (K=8)     | Memory of recent fused observations + goal + state. |
| Action decoder | Cross-attend over T_fused + h_t | Spatial structure reaches the policy head.          |
| Policy Head    | Heteroscedastic ACT, H=8        | Mean + log-variance. No CVAE in v1.                 |
| Goal           | Mandatory relative waypoint     | $g = [p, \sin \alpha, \cos \alpha]$ in robot frame. |
| Execution      | Receding horizon, 1-step        | 2-step only after latency is proven.                |
| Safety         | Independent ROS 2 node          | Higher authority. Uses $\sigma$ and Mahalanobis D.  |

2.3 Initial Dimensions

| Component                     | Value         |
|-------------------------------|---------------|
| DINOv3 backbone               | S+/16, frozen |
| Visual tokens after reduction | 32–64         |
| LiDAR tokens                  | 72–90         |
| Common feature dimension      | 256           |
| Cross-attention heads         | 4             |
| Goal / state embedding        | 32 each       |
| GRU hidden dimension          | 256           |
| Temporal context K            | 8             |
| Action horizon H              | 8             |
| Executed steps per replan     | 1 (initially) |
| Output shape                  | $H \times 4$  |

### 3. Geometry-Aware Fusion & Spatial Path to Actions

#### 3.1 LiDAR Token Representation

For each LiDAR sector/beam  $i$ :  $l_i = [r_i, \sin \theta_i, \cos \theta_i, m_i]$

$r$  = normalized range,  $\theta$  = angle,  $m$  = validity mask. Group nearby beams into 72–90 angular sectors.

#### 3.2 Camera–LiDAR Geometric Mask

Using calibrated intrinsics and extrinsics, classify each LiDAR token as inside or outside the camera FOV. Front/FOV beams attend to compatible visual patches via masked cross-attention. Out-of-FOV beams use a separate LiDAR-only feed-forward path — preventing rear obstacles from being incorrectly associated with front-camera features.

#### 3.3 Gated Fusion

$$\begin{aligned} z_{VL} &= \text{CrossAttn}(T_L^{\text{front}}, T_V, T_V) \\ g &= \sigma(\text{MLP}_g(\text{Pool}(z_{VL}) \parallel \text{Pool}(T_L))) \\ f^{\text{fusion}} &= g \parallel \text{Pool}(z_{VL}) + (1-g) \parallel \text{Pool}(T_L) + \text{Pool}(T_L^{\text{rear}}) \end{aligned}$$

The current-step fused token set  $T_{\text{fused}}$  is retained (not discarded after pooling).  $\text{Pool}(T_{\text{fused}})$  feeds the GRU for temporal memory. The ACT decoder cross-attends over  $T_{\text{fused}}$  with the GRU hidden state  $h_t$  as an additional context token. This ensures geometry-aware spatial structure actually reaches the action outputs.

#### 3.4 Goal & State Integration

$$\begin{aligned} e_{g,t} &= \text{MLP}_g(\rho_t, \sin \alpha_t, \cos \alpha_t) \\ x_t &= [ \text{Pool}(T_{\text{fused}}); e_{g,t}; e_{s,t} ] \\ h_t &= \text{GRU}(x_t, h_{t-1}) \end{aligned}$$

Goal embedding is introduced before the GRU so temporal memory conditions on navigation intent.

### 4. Dual Uncertainty & Training Loss

#### 4.1 Output Parameterization

$$y_{t,h} = [ \mu_{v,h}, \mu_{\omega,h}, \log \sigma_{v,h}^2, \log \sigma_{\omega,h}^2 ]$$

Predicted  $\sigma$  captures **aleatoric** uncertainty (ambiguity/noise among training-like situations). No CVAE is used in v1 — a deterministic chunking head + heteroscedastic NLL is cleaner.

#### 4.2 Mahalanobis OOD (Epistemic) Signal

After data collection, extract pooled DINOv3 feature  $z \in \mathbb{R}^d$  for every training frame. Fit a Gaussian ( $\mu, \Sigma$ ) with shrinkage or diagonal covariance for stability. At deployment compute:

$$D(z) = \sqrt{(z - \mu)^T \Sigma^{-1} (z - \mu)}$$

Large  $D$  means the current visual features lie far from the training distribution (OOD / epistemic signal). Cost is negligible (one matrix-vector product). Both  $\sigma$  and  $D$  are fed to the safety supervisor. This converts frozen SSL features into a free real-time OOD detector for closed-loop safety gating.

#### 4.3 Training Loss with Variance-Collapse Protection

**Warmup phase (first  $N$  epochs):** pure L1 or MSE on action means only.

**Main phase:** heteroscedastic Gaussian NLL

$$L_{\text{NLL}} = \frac{1}{2} \sum_h \sum_{d \in \{v, \omega\}} [ \exp(-s^d) (a^d - \mu^d)^2 + s^d ]$$

where  $s = \log \sigma^2$ . Clip  $s \in [-5, 2]$  early for numerical stability. Warmup prevents the known failure mode of variance collapse on hard samples. Optional: temporal smoothness penalty on consecutive  $\mu$  in the chunk; mild gate regularization.

Full v1 objective:  $L = \lambda_{\text{NLL}} L_{\text{NLL}} + \lambda_{\text{smooth}} L_{\text{smooth}} + \lambda_{\text{gate}} L_{\text{gate}}$

**No CVAE / no KL term in v1.**

### 5. Execution & Safety Supervisor

#### 5.1 Receding-Horizon Execution

Predict  $H=8$  future actions; execute only the first command; obtain new observation; replan. One-step execution is the default for fastest reaction to people, new obstacles, and sensor dropouts. Two-step execution only after latency measurements prove the system can sustain the desired rate.

#### 5.2 Safety Supervisor Hierarchy (Higher Authority than Policy)

$$a_{\text{executed}} = \text{SafetyFilter}(a_{\text{policy}}, z_{\text{LiDAR}}, \sigma, D_{\text{Mahalanobis}}, \Delta t_{\text{sensor}})$$

1. Emergency stop or manual override

2. Sensor or command timeout stop
3. Collision-distance / braking-distance safety stop
4. Velocity and acceleration limiting
5. Uncertainty-based slowdown or stop (uses both  $\sigma$  and  $D$ )
6. Learned-policy command (lowest priority)

### 5.3 Dual-Signal Uncertainty Rule

u combines aleatoric  $\sigma$  and epistemic  $D$ . Example form:

$$v_t^{\text{safe}} = \gamma(\sigma, D) \cdot v_t^{\text{policy}} \text{ with } 0 \leq \gamma \leq 1$$

Thresholds chosen via validation risk–coverage analysis on held-out and novel environments. Key paper figure: collision / intervention rate vs.  $D$ -threshold, showing the OOD gate catches failures that aleatoric  $\sigma$  alone misses.

## 6. Latency Constraint (Critical)

Target control rate: **10 Hz**. Full pipeline (capture → preprocess → inference → safety → publish) must stay under **100 ms** (ideal ≤ 80 ms average). DINOv3-S+/16 runs on the arranged GPU (not on the Raspberry Pi). Measure full round-trip latency early, including any network hop. If marginal, further reduce visual tokens or stay strictly at one-step execution.

## 7. Data Collection Plan

### 7.1 Recording Schema (@ 10 Hz)

| Field       | Source             | Notes                                                  |
|-------------|--------------------|--------------------------------------------------------|
| rgb         | RPi Cam v3         | Train resolution 224×224 (or 320×240 capture)          |
| lidar       | LDS-03             | Ranges + angles; ~5 Hz spin → nearest-scan association |
| action      | Teleop             | linear.x, angular.z                                    |
| robot_state | Odom               | Current $v, \omega$                                    |
| goal        | Planner / operator | Relative $[p, \sin \alpha, \cos \alpha]$ — mandatory   |
| timestamp   | System             | Required for multi-rate association                    |

### 7.2 Protocol Highlights

- Primary store: rosbag2. LeRobot format as exported training view.
- LDS ~5 Hz vs 10 Hz sampling → nearest-timestamp association (expect duplicated scans).
- Report total time, distance, turns, obstacle density, lighting, interventions, near-collisions.
- Train/val/test split by physical environment (not random frames) to avoid leakage.
- 150–200 episodes for first working policy; 300–500 for stronger results.
- Short Gazebo/Isaac sanity check recommended before large-scale real teleoperation.
- Drive smoothly (joystick preferred); deliberately include near-misses and recoveries.

## 8. Baselines & Evaluation

### 8.1 Required Baselines

- Classical Nav2 local navigation
- LiDAR-only learned policy
- RGB-only DINO policy
- Simple RGB–LiDAR concatenation (no cross-attention)
- Cross-attention without gating / without FOV mask
- Full model without temporal memory
- Full model without uncertainty ( $\sigma$  or Mahalanobis)
- Aleatoric-only vs aleatoric + Mahalanobis safety gating

### 8.2 Metrics

**Navigation:** Success Rate, Collision Rate, path efficiency, recovery success.

**Uncertainty:** Failure-prediction AUROC, NLL, calibration error, risk–coverage curves, collision rate at different  $\sigma$  and D thresholds.

**Systems:** End-to-end latency, achieved control frequency, GPU/RAM, parameter count.

## 9. Pilot Implementation Roadmap

| Stage | Focus           | Deliverable                                                                                  |
|-------|-----------------|----------------------------------------------------------------------------------------------|
| 1     | Infrastructure  | ROS 2 recording, camera–LiDAR calibration, goal-conditioned export, latency measurement      |
| 2     | Core model      | Frozen DINO features, LiDAR encoder, geometry-aware fusion, ACT + spatial decoder            |
| 3     | Reliability     | GRU memory, heteroscedastic head (with warmup), Mahalanobis fit, receding-horizon controller |
| 4     | Safety & deploy | Independent safety supervisor ( $\sigma + D$ ), closed-loop real-robot tests                 |
| 5     | Research polish | Baselines, ablations, latency report, risk–coverage and D-threshold figures                  |

## 10. Locked Decisions Summary (v1.1 Final)

| Decision                | Final Choice                                                                         |
|-------------------------|--------------------------------------------------------------------------------------|
| Vision backbone         | Frozen DINOv3-S+/16 only                                                             |
| Visual tokens           | 32–64 after fixed spatial pooling                                                    |
| LiDAR tokens            | 72–90 with angular + validity encoding                                               |
| Fusion                  | Geometry-aware Cross-Attn (FOV-masked) + gated residual + LiDAR-only rear            |
| Spatial path to actions | ACT decoder cross-attends over current $T_{\text{fused}}$ ; GRU $h_t$ as context     |
| Goal conditioning       | Mandatory relative waypoint $[p, \sin \alpha, \cos \alpha]$                          |
| Temporal memory         | GRU, $K = 8$ (pooled path)                                                           |
| Policy                  | Heteroscedastic ACT, $H = 8$ ; no CVAE / no KL in v1                                 |
| Loss                    | L1/MSE warmup $\rightarrow$ full heteroscedastic NLL (+ optional smooth/gate)        |
| Aleatoric uncertainty   | Predicted $\sigma$ from heteroscedastic head                                         |
| Epistemic / OOD         | Mahalanobis distance on pooled frozen DINO features                                  |
| Execution               | Receding horizon, 1-step first                                                       |
| Safety                  | Independent ROS 2 supervisor; priority hierarchy; uses $\sigma$ and $D$              |
| Data                    | rosbag2 primary + LeRobot export; environment-based splits; nearest-scan association |
| Control rate            | 10 Hz target; full pipeline $< 100$ ms                                               |
| Status                  | v1.1 — Final Locked for Pilot Implementation                                         |

**Immediate next actions:** (1) synchronized ROS 2 recording + camera–LiDAR calibration, (2) latency measurement on the arranged GPU, (3) LiDAR-only and concatenation baselines, (4) geometry-aware fusion + spatial ACT training with warmup, (5) Mahalanobis fit on training features, (6) closed-loop deployment behind the safety supervisor.